Yisheng Gong

+1 314-904-6457(US) | yishenggong9437@gmail.com | linkedin | github | website

Education

Washington University in St.Louis — St.Louis, MO
Master of Science in Business Analytics — *Jan 2021 – May 2022*

Fudan University — Shanghai, China
Bachelor of Science in Chemistry — *Sept 2015 – June 2019*

Technical Skills

Languages: Python, SQL, JavaScript
AI/LLM: SFT, RL (GRPO), LoRA & full fine-tuning, distributed training, LLM-as-a-Judge, LLM agents
ML Libraries: PyTorch, CUDA (IPC/streams), torch.compile, Hugging Face (transformers, trl, peft), vLLM, SGLang (contributor), veRL (contributor), miles (contributor), FAISS, LangGraph
Infra & Observability: Kubernetes, KEDA, AWS (EC2, S3, RDS, Lambda), Databricks, Spark, Prometheus, Grafana
Databases: PostgreSQL, Redis, MySQL

Open Source

SGLang/Miles — Dec 2025 -
Contributor

- Cut LoRA weight-update latency by 90% (5s → 0.5s) for colocated training and inference by sharing device weights over CUDA IPC, eliminating host round-trips during RL rollout weight sync.
- Built the distributed LoRA weight-sync path end to end: added the weight-sync API on the SGLang inference side and the distributed weight-sync logic on the miles trainer side, enabling multi-node LoRA post-training.

Professional Experience

Software Engineer, Machine Learning — June 2022 –
Yipitdata, Edison Platform Team — *Mountain View, CA, US*

- Built a production vendor-match system: a dual-encoder + reranker pipeline for routine cases, an LLM agent with KB search for complex queries. Processes 2M+ records/day at 98%+ accuracy, cutting manual workload 75%.
- Optimized the vendor-match GPU inference stack (embedding, reranker, FAISS) for 3x throughput on A10G via micro-batching, CUDA streams, `torch.compile` fusion, and single-GPU data parallelism, with a further 4x from retraining the embedding model on MNR loss. Autoscaled with KEDA/HPA to minimize GPU cost.
- Built a multi-turn SFT + agentic RL (GRPO) post-training pipeline for a 14B LLM with LoRA (negative-sampling traces for SFT; rollout resampling to drop low-variance cases for RL). Lifted ambiguous-case accuracy 60%→90%, beating GPT-5-class baselines
- Migrated a TB-scale ETL pipeline from AWS EMR to Databricks and tuned weekly Spark jobs for 60% faster runtime and 40% lower cost, with automated data-quality checks.
- Optimized core backend infra for Edison Mail via profiling and packet capture: 55% higher throughput, 70% lower cost. Built a Redis client-side cache (-90% I/O) and an event-driven MySQL-binlog → Redis Streams CDC sync.

Software Engineer, Part Time — Jan 2022 – May 2022
Massachusetts Institute of Technology, ChemE, CRIPT Team — *Remote, US*

- Designed and Developed an open-source, semi-automated data ingestion tool based on excel spreadsheet and CRIPT sdk for chemists to upload experiment data to CRIPT database using python. Designed the tool with a compiler-like structure. Assisted early users to upload about 100 thousand data points to database.

Data Engineer Intern — Jan 2021 – May 2021
Bytedance(TikTok), Group of Douyin E-commerce, Commercialization Division — *Beijing, China*

- Developed ETL pipeline with 30+ dependencies to process collected raw data and store in Hbase using HiveSQL. Developed 4 DM tables with different aggregation degrees to satisfy needs. Separated DW data from business logic. Supported the data needs of 200+ colleagues from data science team and operation team.